

Benchmarking **MongoDB and Couchbase** **No-SQL Databases**

Alex Voss

Chris Choi


University of St Andrews

TOP 2 Questions

Should a social scientist *buy*
MORE or **UPGRADE**
computers?

Which **DATABASE(s)**?

Document Oriented Databases




The central concept of a document-oriented database is the notion of a Document

Both



mongoDB and



Couchbase

Uses **JSON** to store these
Documents

Example Documents:

```
{
  FirstName:"Bob",
  Address:"5 Oak St.",
  Hobby:"sailing"
}

{
  FirstName:"Jonathan",
  Address:"15 Wanamassa Point Road",
  Children:[
    {Name:"Michael",Age:10},
    {Name:"Jennifer", Age:8},
    {Name:"Samantha", Age:5},
    {Name:"Elena", Age:2}
  ]
}
```

there are no empty 'fields', and it does not require explicitly stating if other pieces of information are left out

The Database Tests

Database Test

Aggregate/Count:

Most User Mentioned

Most Hashed Tags

Most Shared URLs

Database Test


Aggregate/Count:


Typical questions

Most User Mentioned

Most Hashed Tags

Most Shared URLs





mongoDB



Has two query approaches:

Map Reduce & Aggregation Framework

Map Reduce

For *Aggregation*

Uses **JavaScript**

Aggregation Framework

Similar to Map Reduce


But Uses C++ (Faster!)



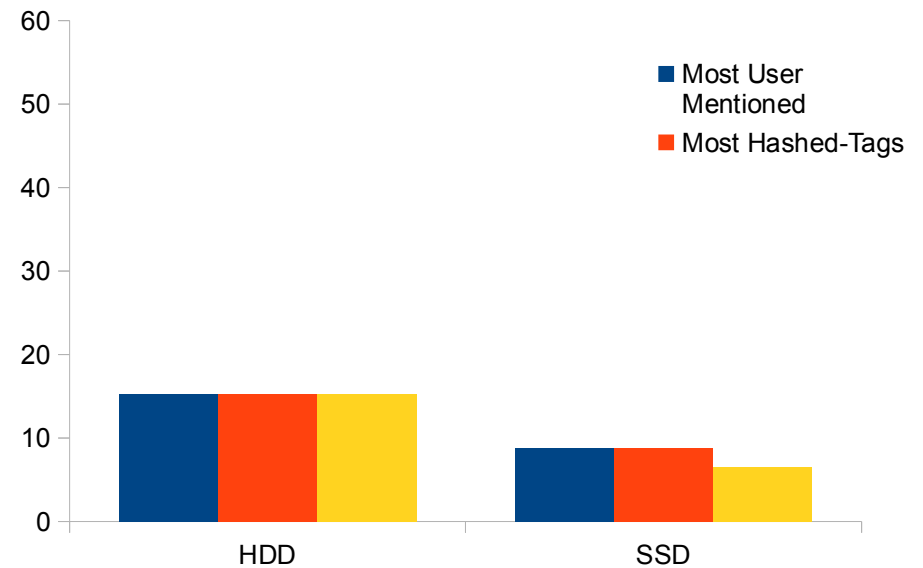
Single Node

SSD vs HDD

Map Reduce



Aggregation Framework





Single Node

Aggregation Framework

is roughly ***2 times faster*** than

Map Reduce

**When we have more than 1
node \sim Scaling**

Generally . . .

Scaling
Vertically



BIGGER
and
FATTER
machine!

Scaling
Horizontally




MORE machines!



Replica-Set

**Master-Slave
Replication**

**Replicates data
from master to
slave**




**Only one node is in
charge (the Master), and
data only travels in one
direction (to the Slave)**



Replica-Set

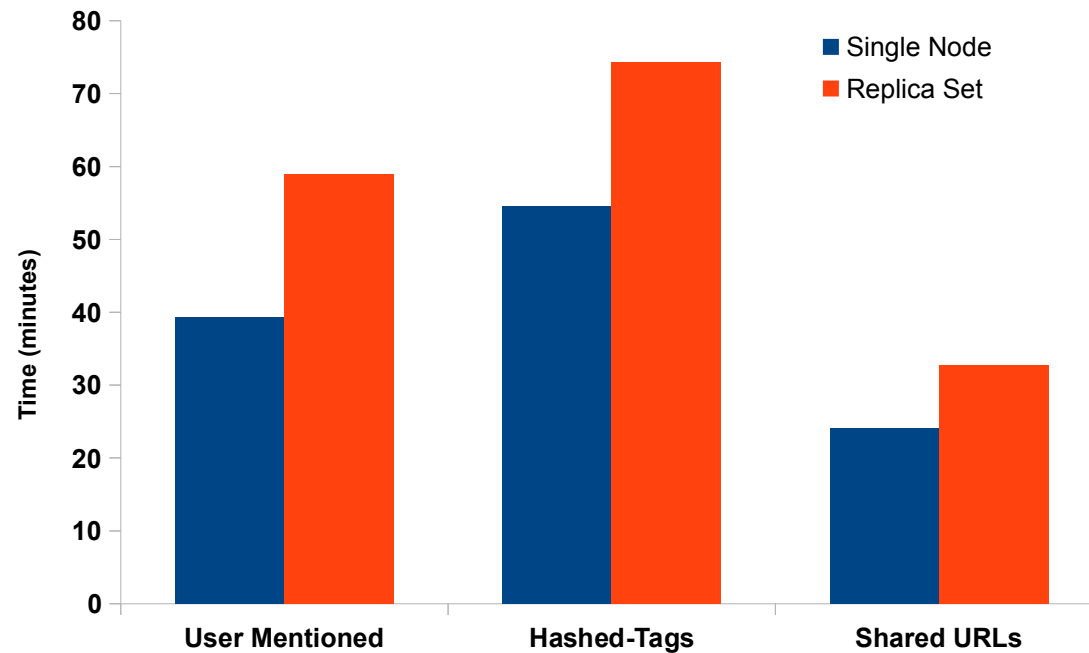
Single Node vs Replica Set: Aggregation Framework





Replica-Set

Single Node vs Replica Set: Map Reduce



***Replicating Data did not improve
query performance***

(MapReduce and Aggregation Framework) . . .

so we tried

Sharding

Sharding

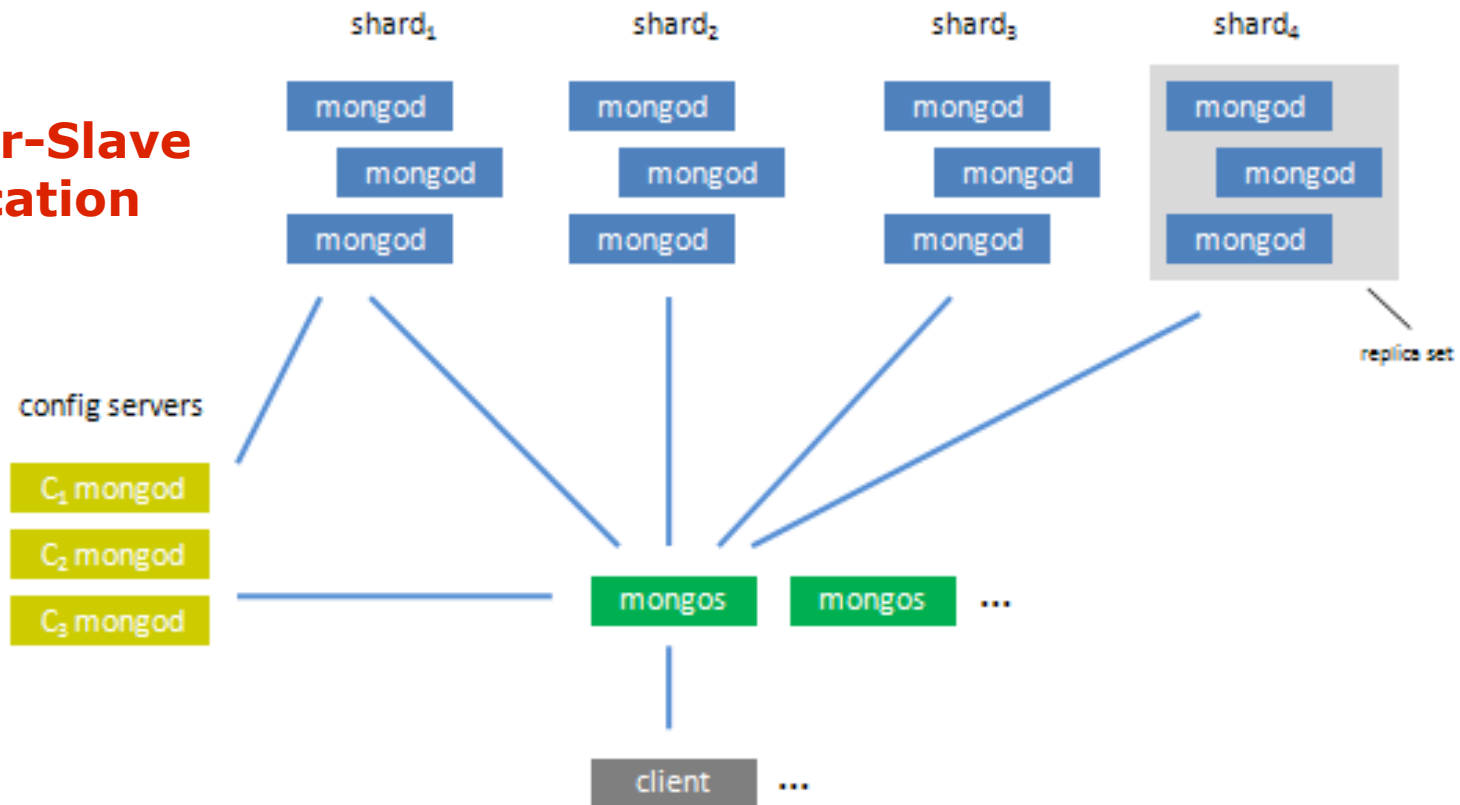


Two or more smaller stacks



uses a **democratic hierarchical** scaling approach

Master-Slave Replication



**But Deployment is an
issue with MongoDB . . .**

If you want to shard . . .



documentation suggests
minimum 11 nodes, to build a
fail-safe *sharded cluster*

2 Shards → 6 nodes (3 each)

2 mongos (load balancer) → 2 nodes

3 mongod (config servers) → 3 nodes

Total: 11 nodes

The *investment* to go from
single node --> **sharded cluster**
is very **HIGH!**



By omitting Fail-Safe features



For our tests, we kind of ***cheated*** and used:

Note: X is variable

~~2~~ **X** Shards → ~~6~~ **2** nodes

~~2~~ **1** mongos (load balancer) → ~~2~~ **0.5** node

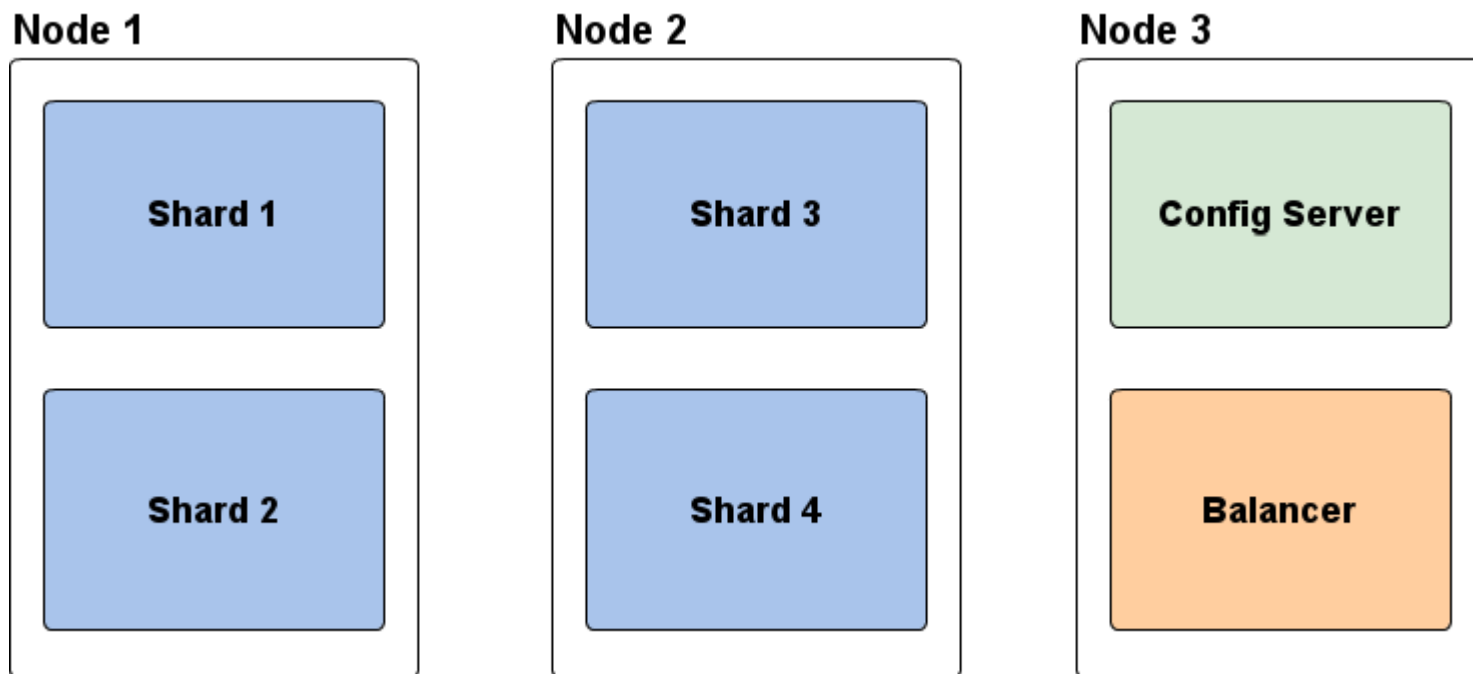
~~3~~ **1** mongod (config servers) → ~~3~~ **0.5** node

Total: ~~11~~ **3** nodes

**NOT RECOMMENDED
IN PRODUCTION**



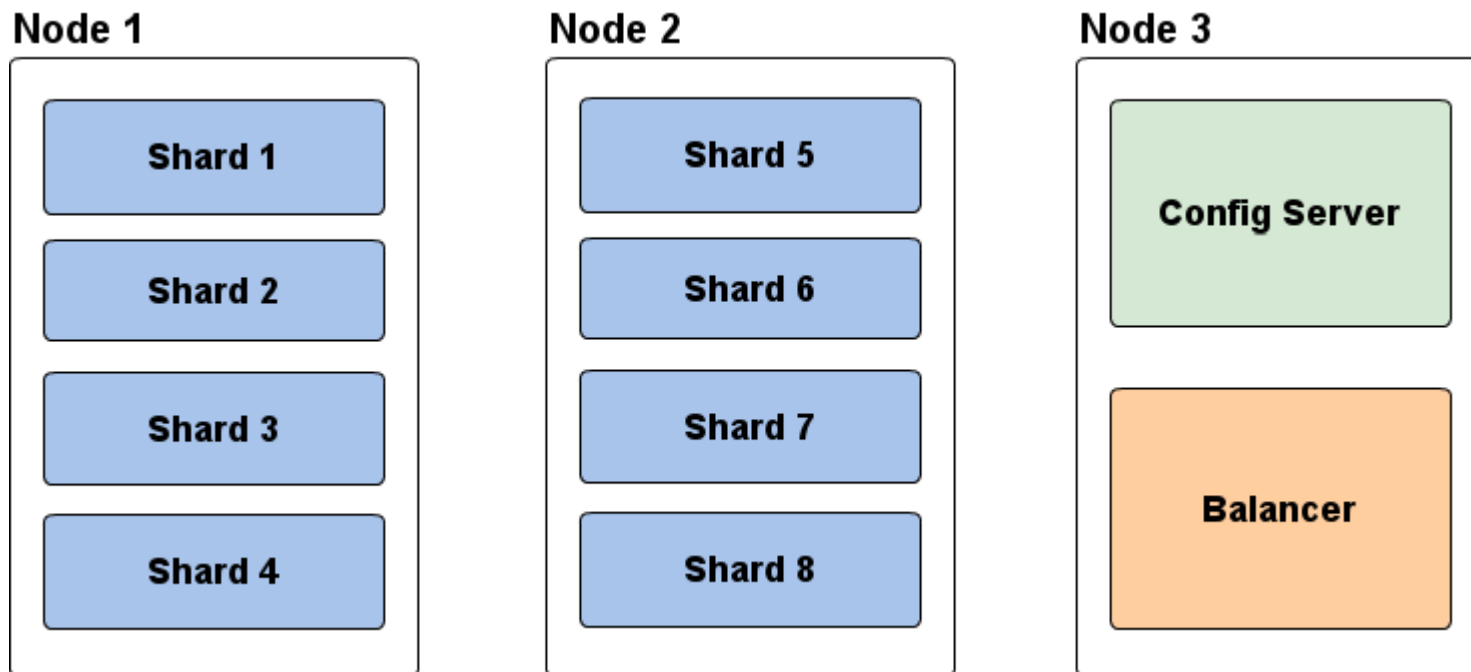
Which looks something like this:



4 Shard configuration



Or this:



8 Shard configuration

**Another problem is with
querying in a sharded
environment . . .**



We ***couldn't*** perform queries using the **Aggregation Framework** in the shards, because it was limited by:

- Output at every stage of the aggregation can only contain **16MB**
- **>10%** RAM usage

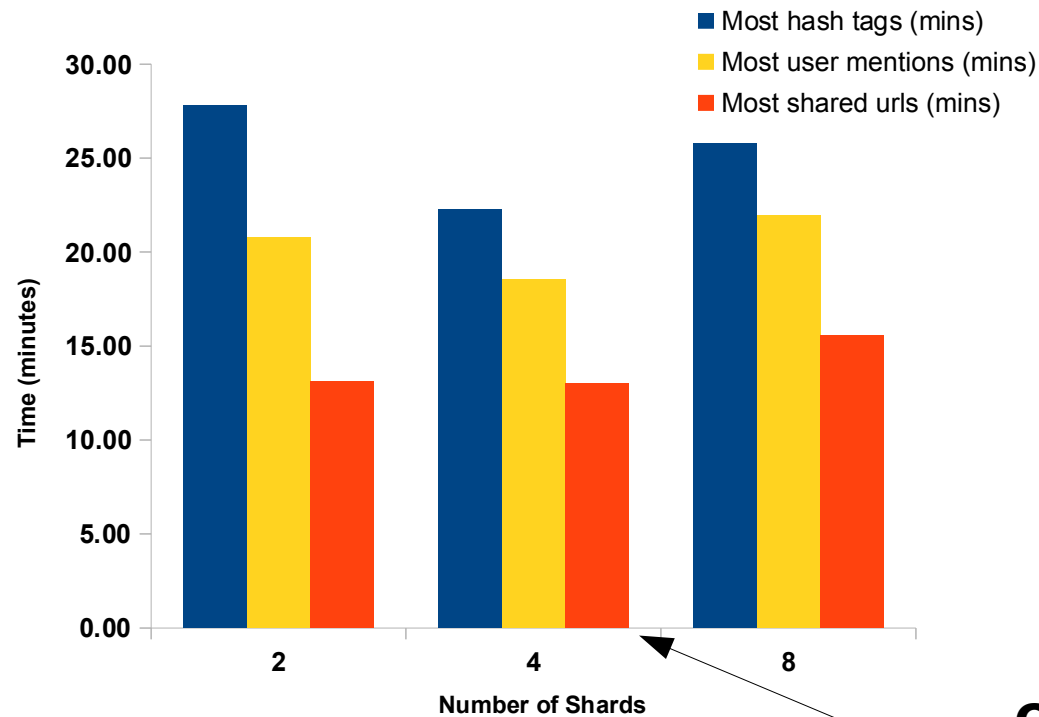
MongoDB would fail if the limits were reached



4 Cores
8GB RAM

Sharded Cluster

Map Reduce: On 3 modest machines



Optimal

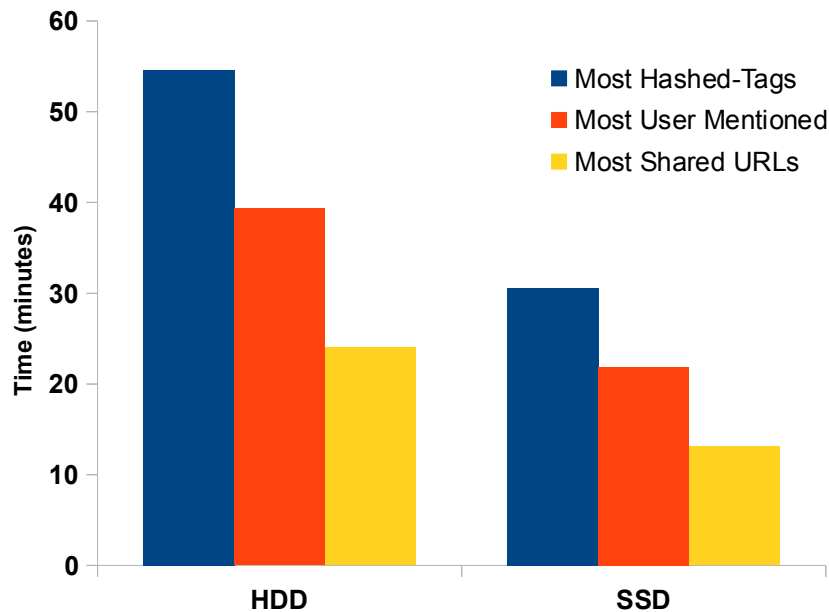


4 Cores
8GB RAM


Sharded Cluster

Map Reduce: On 3 modest machines

Single Node



Three Nodes






4 Cores
8GB RAM

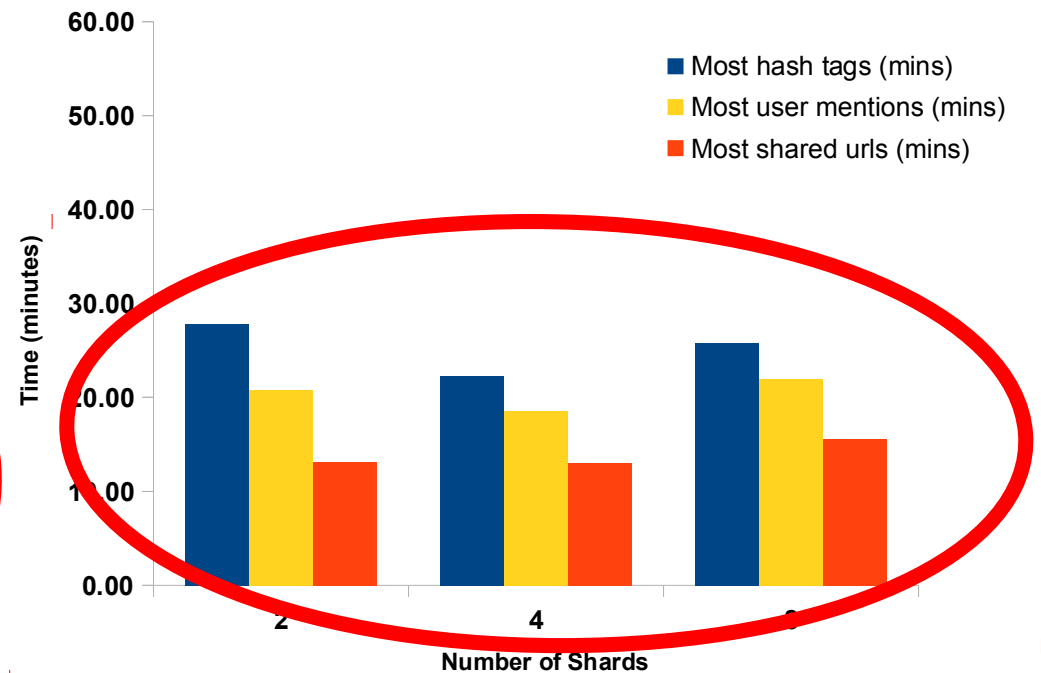
Sharded Cluster

Map Reduce: On 3 modest machines

Single Node



Three Nodes





**Performance of Single Node (SSD)
is very comparable to 3 nodes (HDD)
of different number of shards!**

**(Single Node with SSD is
2-3 minutes slower than 3
nodes . . . [40 GB corpus])**

**Perhaps its worth purchasing a good
SSD over *investing more nodes* to
improve aggregate performance . . .**



What happens when we scale vertically
with **ONE BIG BOX?**


we get --->




32 Cores
128GB RAM

Sharded Cluster


Map Reduce: On a Big Box




Optimal



CouchBase



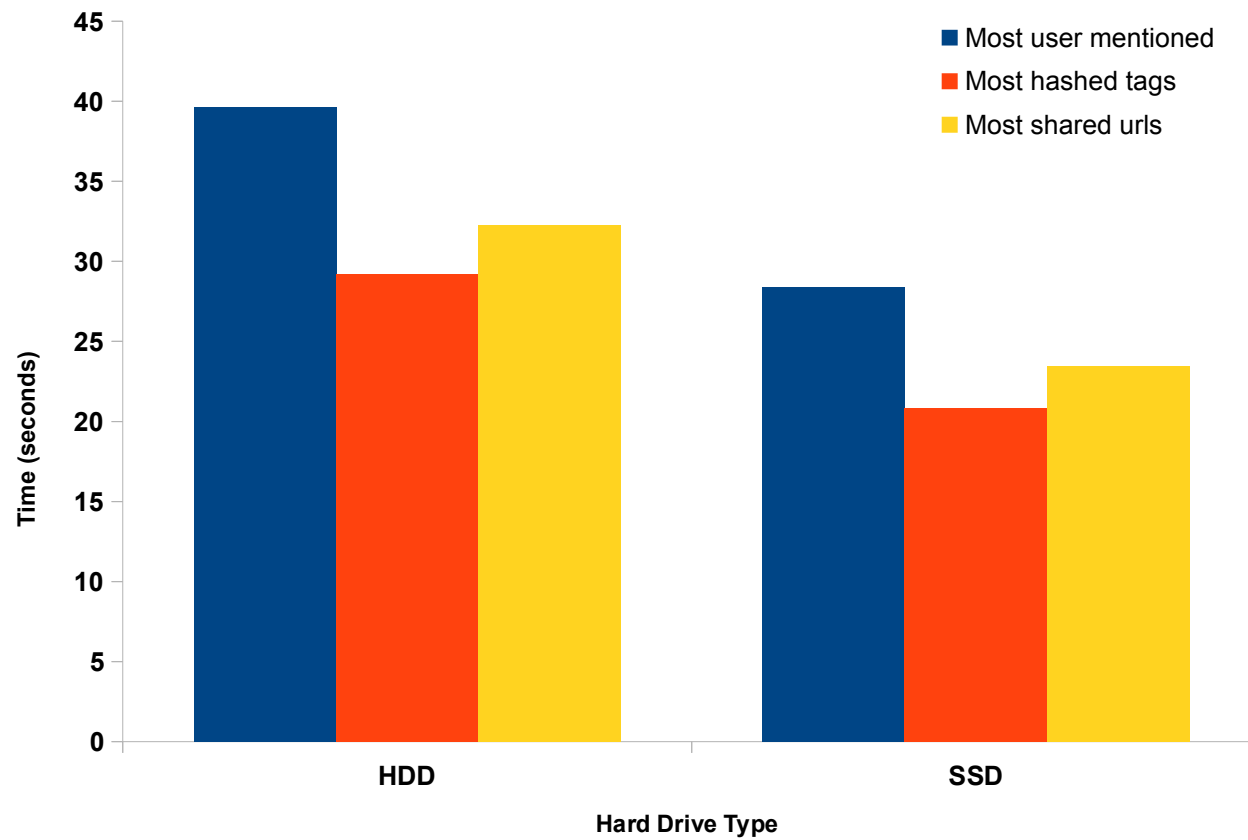
uses a **communistic**
scaling approach



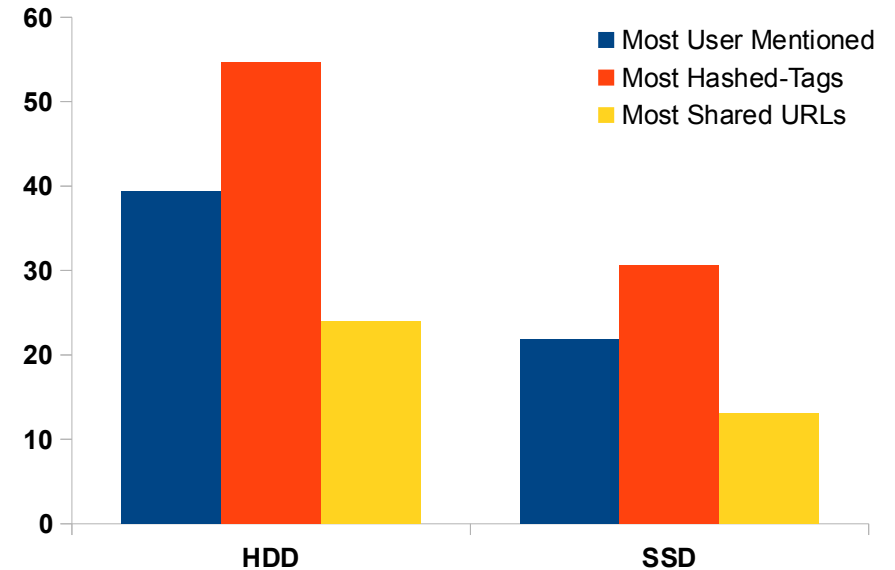

Master-Master Replication




Single Node



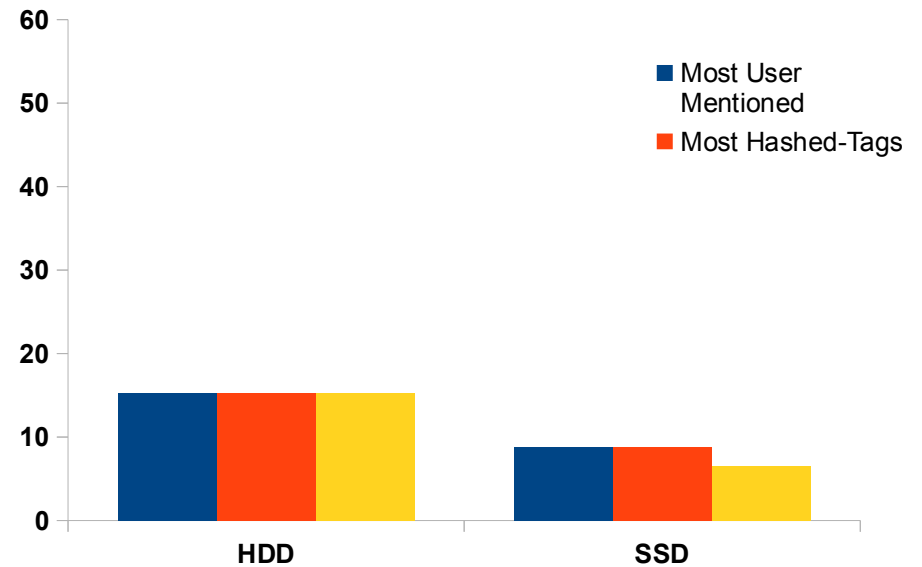
Single Node VS



Single Node VS



Aggregation Framework



On a **SINGLE NODE** , while

CouchBase
is ***FASTER*** than
MongoDB's MapReduce

MongoDB's Aggregation Framework
is ***FASTER*** than
CouchBase

**That said adding more
Couchbase nodes was
disappointing. . .**

In short . . .




**1
Node**

vs

**2
Nodes**

vs

**3
Nodes**



***Adding more nodes did not
improve
query performance . . .***

**Other users have found similar issues,
Couchbase Support says:**

**“One known issue with the current developer
preview of Couchbase is that it *slows down*
on *smaller clusters*. The views are optimal
on *very large clusters*.”** [1]

But they did not specify what is “very large”

[1] [Couchbase performance - real tests - it's horrible - or am I doing somethin wrong?](#)

That was

November 2011 . . .

Which one?



The obvious answer is



From a performance standpoint.

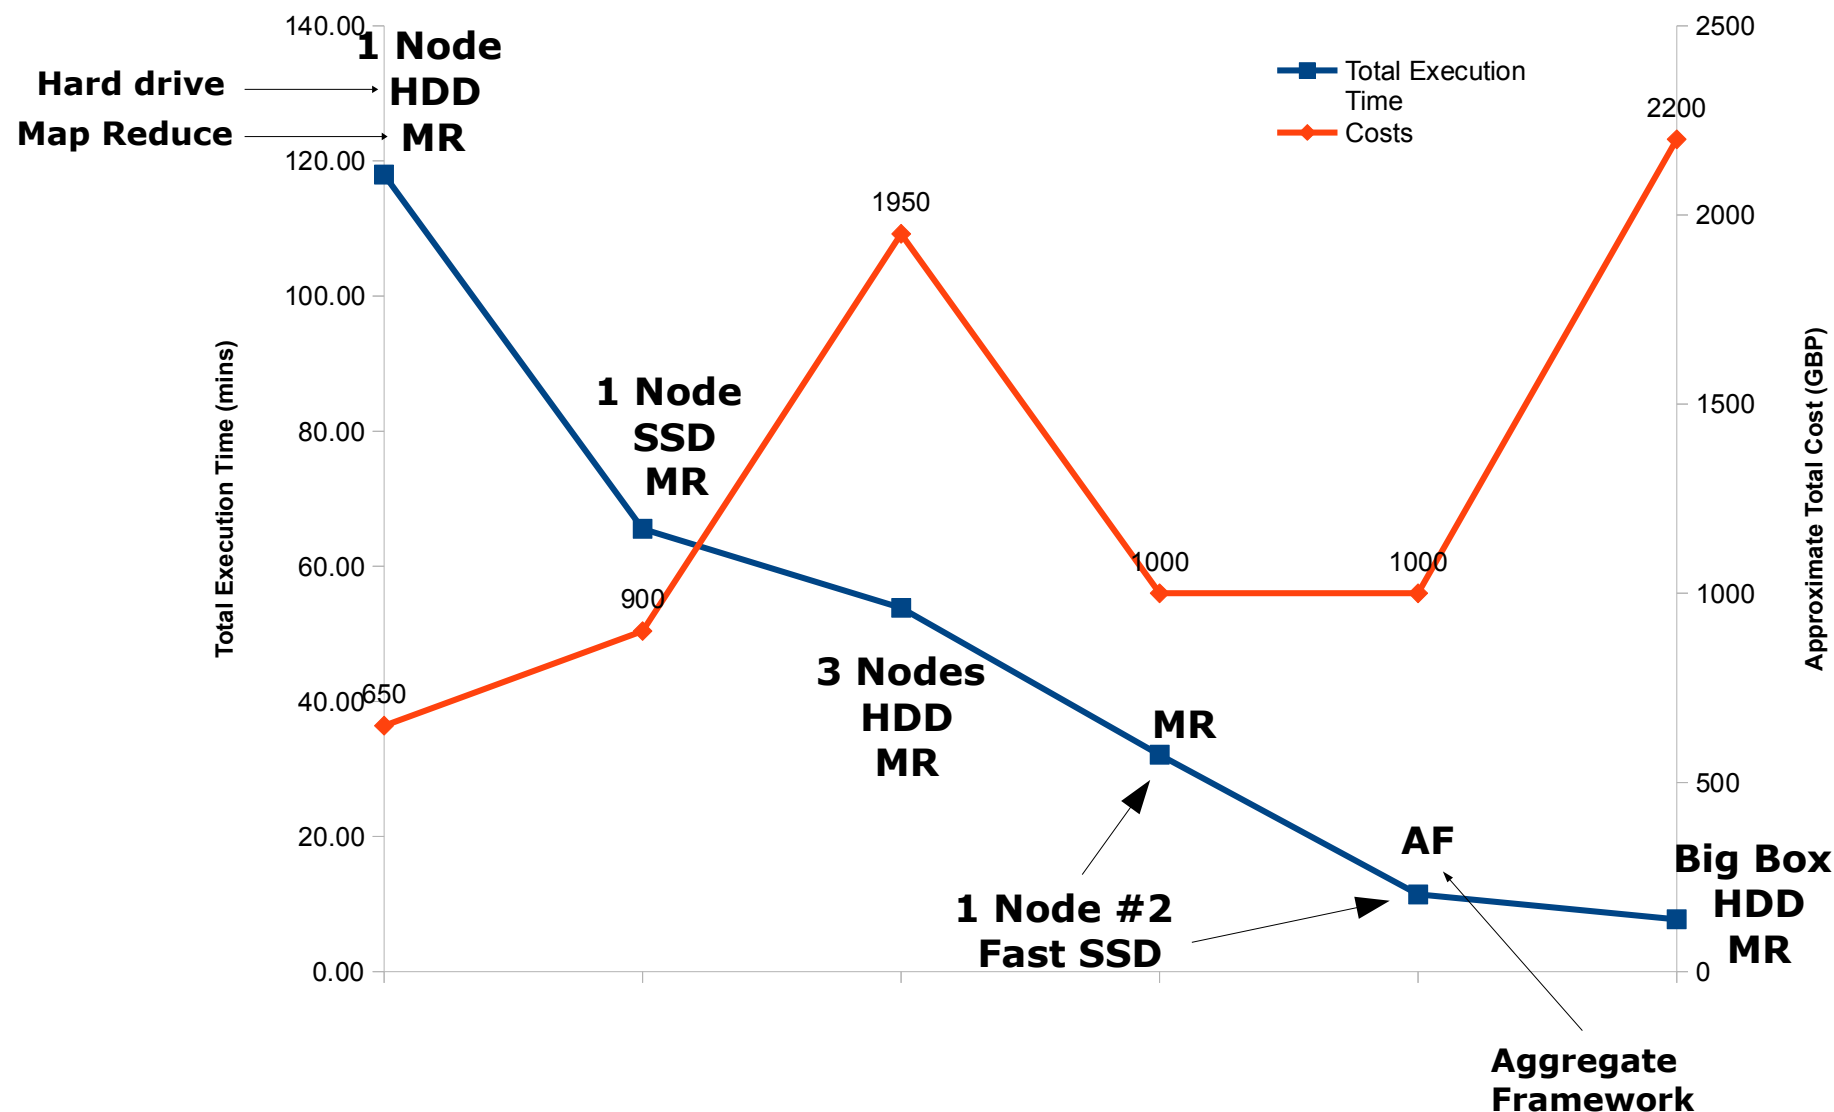
Simple flat file aggregation with *Java*
took approx 60 mins

On the same machine, MongoDB took
approx

30 mins (MapReduce)

10 mins (Aggregation Framework)

Deployment Options



But!



mongoDB Has problems with:

- Full Text Search
- Social Network Graphs
- Aggregation



mongoDB's Full Text Search
is **slow**

1 simple query on a single
node, 44GB corpus took
approx

10 mins!

Solution:



A **Fast** Full Text Search Engine

The same search took only a fraction of a second


To create a
Social
Network
Graph with



mongoDB

Can be ***Complex***,
since you query multiple
times to obtain
relationships

Solution:



A ***Popular*** Graph Database

Where you can obtain a
network graph
with **VERY FEW** queries

Example:

```
START john=node:node_auto_index(name = 'John')  
MATCH john-[:friend]->()-[:friend]->fof  
RETURN john, fof
```

**The above query
fetches all nodes who
are friends of friends
of “Johns”**

With Aggregation



***Is perhaps* not the best,
we have yet to play with**



Or even



An **SQL** language to
build MapReduce
queries

With




Moral of the story

**Use the right tools
for the right job**

Future work

We aim to experiment with:



**With possibility in integrating all
with**

